

PRINCE SAVSAVIYA

+1-951-377-3380 | princesavsaviya2023.learning@gmail.com | linkedin.com/in/prince-savsaviya | github.com/princesavsaviya | princesavsaviya.github.io/profile/

Professional Summary

Machine Learning Engineer with experience building scalable ML systems across retrieval, time-series forecasting, physics-informed modeling, and cloud deployment. Strong background in **model development, feature engineering, hyperparameter tuning, and low-latency inference pipelines** using **PyTorch, TensorFlow, Scikit-learn, AWS, and Kubernetes**. Proven impact improving search precision by **25%**, reducing latency by **40%**, and accelerating experimentation through automated ML workflows.

Education

University of California, Riverside M.S. Computer Science, GPA 3.56/4.0	Riverside, CA	March 2026
Adani Institute of Infrastructure Engineering, GTU B.E. Information & Communication Technology, GPA 3.6/4.0	Ahmedabad, IN	May 2024

Core Technical Skills

Machine Learning:	Deep Learning, PyTorch, TensorFlow, Scikit-learn, BERT, LLMs (Hugging Face), Computer Vision, Physics-Informed ML, Time-Series Forecasting, Hyperparameter Tuning
Backend & APIs:	Django, FastAPI, REST APIs, Microservices
MLOps & Cloud:	Docker, Kubernetes, AWS (EKS, EC2, S3), GitHub Actions (CI/CD), MLflow, Optuna
Languages:	Python, C++, SQL, Bash, CUDA, C, Java
Big Data & Databases:	PySpark, Kafka, Hadoop, ETL Pipelines, PostgreSQL, Elasticsearch, FAISS
Data & Analytics:	Pandas, NumPy, Data Analysis, Data Cleaning, Data Visualization
Systems:	MPI, PETSc, DofinX

Experience

Graduate Researcher — University of California, Riverside	Riverside, CA	May 2025 – March 2026
- Built ML-ready data pipelines from large-scale multiphysics simulations for surrogate modeling and physics-informed learning tasks. - Designed and evaluated physics-informed ML models, including Hamiltonian Neural Networks, to reduce solver iterations and improve convergence efficiency. - Engineered distributed experimentation workflows using Python, Bash, and MPI, improving throughput by 2×. - Developed automated benchmarking and regression-testing pipelines, reducing setup time for new cases by 80%.		
Research Intern — Adani University	Ahmedabad, IN	May 2023 – Dec 2023
- Built a physics-informed ML framework combining PDE ageing models with ensemble regressors (SVR, RF, GBR), cutting MSE 98% and RMSE 86% vs. data-only baseline. - Ingested 200,000+ battery-cycle records in Python; automated feature engineering and Optuna hyperparameter tuning, reducing experiment turnaround 30%. Publication: <i>Journal of Power Sources</i>, Vol. 627, 2025 (IF ≈ 7.9) — DOI 10.1016/j.jpowsour.2024.235771.		

Selected Projects

RetrieveX — Semantic Search Engine	Django · BERT · FAISS · AWS	Jan 2025 – Apr 2025
- Architected a hybrid retrieval system merging BM25 keyword search with BERT+FAISS vector embeddings, improving top-3 precision 25% on 150k medical documents. - Engineered an asynchronous inference pipeline using Django and REST APIs, optimizing query latency to <200 ms (P_{95}).		
AI-Driven Infrastructure Autoscaler	LSTM · Kubernetes · Helm · CI/CD	Jan 2025 – Apr 2025
- Designed a time-series forecasting service using LSTMs to predict cluster load, automating replica scaling and cutting cloud over-provisioning costs by 18%. - Implemented MLOps best practices with automated model retraining triggers and CI/CD deployment to AWS EKS; sustained 99.9% SLA under 4× load spikes.		
Risk-Analytics Microservice	FastAPI · Pandas · Docker · AWS	Sep 2024 – Jan 2025
- Designed real-time analytics service ingesting multi-asset data with vectorized pipelines; delivered metrics in <150 ms. - Automated deployment with GitHub Actions → AWS ECR → EC2; maintained 99.9% uptime with Prometheus monitoring.		